



Monitor a storage system

ONTAP 9

NetApp
May 31, 2023

Table of Contents

- Monitor a storage system 1
 - Use AutoSupport and Active IQ Digital Advisor 1
 - Manage AutoSupport settings with System Manager 1
 - Manage AutoSupport with the CLI 3
- Monitor the health of your system 29
- Display environmental information 42

Monitor a storage system

Use AutoSupport and Active IQ Digital Advisor

The AutoSupport component of ONTAP collects telemetry and sends it for analysis. Active IQ Digital Advisor analyzes the data from AutoSupport and provides proactive care and optimization. Using artificial intelligence, Active IQ can identify potential problems and help you resolve them before they impact your business.

Active IQ enables you to optimize your data infrastructure across your global hybrid cloud by delivering actionable predictive analytics and proactive support through a cloud-based portal and mobile app. Data-driven insights and recommendations from Active IQ are available to all NetApp customers with an active SupportEdge contract (features vary by product and support tier).

Here are some things you can do with Active IQ:

- Plan upgrades. Active IQ identifies issues in your environment that can be resolved by upgrading to a newer version of ONTAP and the Upgrade Advisor component helps you plan for a successful upgrade.
- View system wellness. Your Active IQ dashboard reports any issues with wellness and helps you correct those issues. Monitor system capacity to make sure you never run out of storage space. View support cases for your system.
- Manage performance. Active IQ shows system performance over a longer period than you can see in System Manager. Identify configuration and system issues that are impacting your performance.
- Maximize efficiency. View storage efficiency metrics and identify ways to store more data in less space.
- View inventory and configuration. Active IQ displays complete inventory and software and hardware configuration information. See when service contracts are expiring and renew them to ensure you remain supported.

Related information

[NetApp Documentation: Active IQ Digital Advisor](#)

[Launch Active IQ](#)

[SupportEdge Services](#)

Manage AutoSupport settings with System Manager

You can use System Manager to view and edit the settings for your AutoSupport account.

You can perform the following procedures:

- [View AutoSupport settings](#)
- [Generate and send AutoSupport data](#)
- [Test the connection to AutoSupport](#)
- [Enable or disable AutoSupport](#)
- [Suppress the generation of support cases](#)
- [Resume the generation of support cases](#)

- [Edit AutoSupport settings](#)

View AutoSupport settings

You can use System Manager to view the settings for your AutoSupport account.

Steps

1. In System Manager, click **Cluster > Settings**.

In the **AutoSupport** section, the following information is displayed:

- Status
- Transport protocol
- Proxy server
- From email address

2. In the **AutoSupport** section, click , then click **More Options**.

Additional information is displayed about the AutoSupport connection and email settings. Also, the transfer history of messages is listed.

Generate and send AutoSupport data

In System Manager, you can initiate the generation of AutoSupport messages and choose from which cluster node or nodes the data is collected.

Steps

1. In System Manager, click **Cluster > Settings**.
2. In the **AutoSupport** section, click , then click **Generate and Send**.
3. Enter a subject.
4. Click the check box under **Collect Data From** to specify the nodes from which to collect the data.

Test the connection to AutoSupport

From System Manager, you can send a test message to verify the connection to AutoSupport.

Steps

1. In System Manager, click **Cluster > Settings**.
2. In the **AutoSupport** section, click , then click **Test Connectivity**.
3. Enter a subject for the message.

Enable or disable AutoSupport

In System Manager, you can disable the ability of AutoSupport to monitor the health of your storage system and send you notification messages. You can enable AutoSupport again after it has been disabled.

Steps

1. Click **Cluster > Settings**.
2. In the **AutoSupport** section, click , then click **Disable**.

3. If want to enable AutoSupport again, in the **AutoSupport** section, click , then click **Enable**.

Suppress the generation of support cases

Beginning with ONTAP 9.10.1, you can use System Manager to send a request to AutoSupport to suppress the generation of support cases.

About this task

To suppress the generation of support cases, you specify the nodes and number of hours for which you want the suppression to occur.

Suppressing support cases can be especially helpful if you do not want AutoSupport to create automated cases while you are performing maintenance on your systems.

Steps

1. Click **Cluster > Settings**.
2. In the **AutoSupport** section, click , then click **Suppress Support Case Generation**.
3. Enter the number of hours that you want the suppression to occur.
4. Select the nodes for which you want the suppression to occur.

Resume the generation of support cases

Beginning with ONTAP 9.10.1, you can use System Manager to resume the generation of support cases from AutoSupport if it has been suppressed.

Steps

1. Click **Cluster > Settings**.
2. In the **AutoSupport** section, click , then click **Resume Support Case Generation**.
3. Select the nodes for which you want the generation to resume.

Edit AutoSupport settings

You can use System Manager to modify the connection and email settings for your AutoSupport account.

Steps

1. Click **Cluster > Settings**.
2. In the **AutoSupport** section, click , then click **More Options**.
3. In the **Connections** section or the **Email** section, click  **Edit** to modify the setting for either section.

Manage AutoSupport with the CLI

Manage AutoSupport overview

AutoSupport is a mechanism that proactively monitors the health of your system and automatically sends messages to NetApp technical support, your internal support organization, and a support partner. Although AutoSupport messages to technical support are enabled by default, you must set the correct options and have a valid mail host to have messages sent to your internal support organization.

Only the cluster administrator can perform AutoSupport management. The storage virtual machine (SVM) administrator has no access to AutoSupport.

AutoSupport is enabled by default when you configure your storage system for the first time. AutoSupport begins sending messages to technical support 24 hours after AutoSupport is enabled. You can shorten the 24-hour period by upgrading or reverting the system, modifying the AutoSupport configuration, or changing the system time to be something other than a 24-hour period.



You can disable AutoSupport at any time, but you should leave it enabled. Enabling AutoSupport can significantly help speed problem determination and resolution should a problem occur on your storage system. By default, the system collects AutoSupport information and stores it locally, even if you disable AutoSupport.

For more information about AutoSupport, see the NetApp Support Site.

Related information

- [NetApp Support](#)
- [Learn more about the AutoSupport commands in the ONTAP CLI](#)

When and where AutoSupport messages are sent

AutoSupport sends messages to different recipients, depending on the type of message. Learning when and where AutoSupport sends messages can help you understand messages that you receive through email or view on the Active IQ (formerly known as My AutoSupport) web site.

Unless specified otherwise, settings in the following tables are parameters of the `system node autosupport modify` command.

Event-triggered messages

When events occur on the system that require corrective action, AutoSupport automatically sends an event-triggered message.

| When the message is sent | Where the message is sent |
|--|---|
| AutoSupport responds to a trigger event in the EMS | Addresses specified in <code>-to</code> and <code>-noteto</code> . (Only critical, service-affecting events are sent.) Addresses specified in <code>-partner-address</code> Technical support, if <code>-support</code> is set to <code>enable</code> |

Scheduled messages

AutoSupport automatically sends several messages on a regular schedule.

| When the message is sent | Where the message is sent |
|---|--|
| Daily (by default, sent between 12:00 a.m. and 1:00 a.m. as a log message) | Addresses specified in <code>-partner-address</code> Technical support, if <code>-support</code> is set to <code>enable</code> |
| Daily (by default, sent between 12:00 a.m. and 1:00 a.m. as a performance message), if the <code>-perf</code> parameter is set to <code>true</code> | Addresses specified in <code>-partner-address`</code> Technical support, if <code>-support</code> is set to <code>enable</code> |
| Weekly (by default, sent Sunday between 12:00 a.m. and 1:00 a.m.) | Addresses specified in <code>-partner-address</code> Technical support, if <code>-support</code> is set to <code>enable</code> |

Manually triggered messages

You can manually initiate or resend an AutoSupport message.

| When the message is sent | Where the message is sent |
|---|---|
| You manually initiate a message using the <code>system node autosupport invoke</code> command | <p>If a URI is specified using the <code>-uri</code> parameter in the <code>system node autosupport invoke</code> command, the message is sent to that URI.</p> <p>If <code>-uri</code> is omitted, the message is sent to the addresses specified in <code>-to</code> and <code>-partner-address</code>. The message is also sent to technical support if <code>-support</code> is set to <code>enable</code>.</p> |
| You manually initiate a message using the <code>system node autosupport invoke-core-upload</code> command | <p>If a URI is specified using the <code>-uri</code> parameter in the <code>system node autosupport invoke-core-upload</code> command, the message is sent to that URI, and the core dump file is uploaded to the URI.</p> <p>If <code>-uri</code> is omitted in the <code>system node autosupport invoke-core-upload</code> command, the message is sent to technical support, and the core dump file is uploaded to the technical support site.</p> <p>Both scenarios require that <code>-support</code> is set to <code>enable</code> and <code>-transport</code> is set to <code>https</code> or <code>http</code>.</p> <p>Due to the large size of core dump files, the message is not sent to the addresses specified in the <code>-to</code> and <code>-partner-addresses</code> parameters.</p> |

| When the message is sent | Where the message is sent |
|---|---|
| You manually initiate a message using the <code>system node autosupport invoke-performance-archive</code> command | <p>If a URI is specified using the <code>-uri</code> parameter in the <code>system node autosupport invoke-performance-archive</code> command, the message is sent to that URI, and the performance archive file is uploaded to the URI.</p> <p>If <code>-uri</code> is omitted in the <code>system node autosupport invoke-performance-archive</code>, the message is sent to technical support, and the performance archive file is uploaded to the technical support site.</p> <p>Both scenarios require that <code>-support</code> is set to <code>enable</code> and <code>-transport</code> is set to <code>https</code> or <code>http</code>.</p> <p>Due to the large size of performance archive files, the message is not sent to the addresses specified in the <code>-to</code> and <code>-partner-addresses</code> parameters.</p> |
| You manually resend a past message using the <code>system node autosupport history retransmit</code> command | Only to the URI that you specify in the <code>-uri</code> parameter of the <code>system node autosupport history retransmit</code> command |

Messages triggered by technical support

Technical support can request messages from AutoSupport using the AutoSupport OnDemand feature.

| When the message is sent | Where the message is sent |
|--|---|
| When AutoSupport obtains delivery instructions to generate new AutoSupport messages | <p>Addresses specified in <code>-partner-address</code></p> <p>Technical support, if <code>-support</code> is set to <code>enable</code> and <code>-transport</code> is set to <code>https</code></p> |
| When AutoSupport obtains delivery instructions to resend past AutoSupport messages | Technical support, if <code>-support</code> is set to <code>enable</code> and <code>-transport</code> is set to <code>https</code> |
| When AutoSupport obtains delivery instructions to generate new AutoSupport messages that upload core dump or performance archive files | Technical support, if <code>-support</code> is set to <code>enable</code> and <code>-transport</code> is set to <code>https</code> . The core dump or performance archive file is uploaded to the technical support site. |

How AutoSupport creates and sends event-triggered messages

AutoSupport creates event-triggered AutoSupport messages when the EMS processes a trigger event. An event-triggered AutoSupport message alerts recipients to problems that require corrective action and contains only information that is relevant to the problem. You

can customize what content to include and who receives the messages.

AutoSupport uses the following process to create and send event-triggered AutoSupport messages:

1. When the EMS processes a trigger event, EMS sends AutoSupport a request.

A trigger event is an EMS event with an AutoSupport destination and a name that begins with a `callhome.` prefix.

2. AutoSupport creates an event-triggered AutoSupport message.

AutoSupport collects basic and troubleshooting information from subsystems that are associated with the trigger to create a message that includes only information that is relevant to the trigger event.

A default set of subsystems is associated with each trigger. However, you can choose to associate additional subsystems with a trigger by using the `system node autosupport trigger modify` command.

3. AutoSupport sends the event-triggered AutoSupport message to the recipients defined by the `system node autosupport modify` command with the `-to`, `-noteto`, `-partner-address`, and `-support` parameters.

You can enable and disable delivery of AutoSupport messages for specific triggers by using the `system node autosupport trigger modify` command with the `-to` and `-noteto` parameters.

Example of data sent for a specific event

The `storage shelf PSU failed` EMS event triggers a message that contains basic data from the Mandatory, Log Files, Storage, RAID, HA, Platform, and Networking subsystems and troubleshooting data from the Mandatory, Log Files, and Storage subsystems.

You decide that you want to include data about NFS in any AutoSupport messages sent in response to a future `storage shelf PSU failed` event. You enter the following command to enable troubleshooting-level data for NFS for the `callhome.shlf.ps.fault` event:

```
cluster1::\>
system node autosupport trigger modify -node node1 -autosupport
-message shlf.ps.fault -troubleshooting-additional nfs
```

Note that the `callhome.` prefix is dropped from the `callhome.shlf.ps.fault` event when you use the `system node autosupport trigger` commands, or when referenced by AutoSupport and EMS events in the CLI.

Types of AutoSupport messages and their content

AutoSupport messages contain status information about supported subsystems. Learning what AutoSupport messages contain can help you interpret or respond to messages that you receive in email or view on the Active IQ (formerly known as My AutoSupport) web site.

| Type of message | Type of data the message contains |
|--|--|
| Event-triggered | Files containing context-sensitive data about the specific subsystem where the event occurred |
| Daily | Log files |
| Performance | Performance data sampled during the previous 24 hours |
| Weekly | Configuration and status data |
| Triggered by the <code>system node autosupport invoke</code> command | <p>Depends on the value specified in the <code>-type</code> parameter:</p> <ul style="list-style-type: none"> • <code>test</code> sends a user-triggered message with some basic data. <p>This message also triggers an automated email response from technical support to any specified email addresses, using the <code>-to</code> option, so that you can confirm that AutoSupport messages are being received.</p> <ul style="list-style-type: none"> • <code>performance</code> sends performance data. • <code>all</code> sends a user-triggered message with a complete set of data similar to the weekly message, including troubleshooting data from each subsystem. <p>Technical support typically requests this message.</p> |
| Triggered by the <code>system node autosupport invoke-core-upload</code> command | Core dump files for a node |
| Triggered by the <code>system node autosupport invoke-performance-archive</code> command | Performance archive files for a specified period of time |

| Type of message | Type of data the message contains |
|-----------------------------------|--|
| Triggered by AutoSupport OnDemand | <p>AutoSupport OnDemand can request new messages or past messages:</p> <ul style="list-style-type: none"> • New messages, depending on the type of AutoSupport collection, can be <code>test</code>, <code>all</code>, or <code>performance</code>. • Past messages depend on the type of message that is resent. <p>AutoSupport OnDemand can request new messages that upload the following files to the NetApp Support Site at mysupport.netapp.com:</p> <ul style="list-style-type: none"> • Core dump • Performance archive |

What AutoSupport subsystems are

Each subsystem provides basic and troubleshooting information that AutoSupport uses for its messages. Each subsystem is also associated with trigger events that allow AutoSupport to collect from subsystems only information that is relevant to the trigger event.

AutoSupport collects context-sensitive content. You can view information about subsystems and trigger events by using the `system node autosupport trigger show` command.

AutoSupport size and time budgets

AutoSupport collects information, organized by subsystem, and enforces a size and time budget on content for each subsystem. As storage systems grow, AutoSupport budgets provide control over the AutoSupport payload, which in turn provides scalable delivery of AutoSupport data.

AutoSupport stops collecting information and truncates the AutoSupport content if the subsystem content exceeds its size or time budget. If the content cannot be truncated easily (for example, binary files), AutoSupport omits the content.

You should modify the default size and time budgets only if asked to do so by NetApp Support. You can also review the default size and time budgets of the subsystems by using the `autosupport manifest show` command.

Files sent in event-triggered AutoSupport messages

Event-triggered AutoSupport messages only contain basic and troubleshooting information from subsystems that are associated with the event that caused AutoSupport to generate the message. The specific data helps NetApp support and support partners troubleshoot the problem.

AutoSupport uses the following criteria to control content in event-triggered AutoSupport messages:

- Which subsystems are included

Data is grouped into subsystems, including common subsystems, such as Log Files, and specific subsystems, such as RAID. Each event triggers a message that contains only the data from specific subsystems.

- The detail level of each included subsystem

Data for each included subsystem is provided at a basic or troubleshooting level.

You can view all possible events and determine which subsystems are included in messages about each event using the `system node autosupport trigger show` command with the `-instance` parameter.

In addition to the subsystems that are included by default for each event, you can add additional subsystems at either a basic or a troubleshooting level using the `system node autosupport trigger modify` command.

Log files sent in AutoSupport messages

AutoSupport messages can contain several key log files that enable technical support staff to review recent system activity.

All types of AutoSupport messages might include the following log files when the Log Files subsystem is enabled:

| Log file | Amount of data included from the file |
|---|---|
| <ul style="list-style-type: none">• Log files from the <code>/mroot/etc/log/mlog/</code> directory• The MESSAGES log file | Only new lines added to the logs since the last AutoSupport message up to a specified maximum. This ensures that AutoSupport messages have unique, relevant—not overlapping—data. (Log files from partners are the exception; for partners, the maximum allowed data is included.) |
| <ul style="list-style-type: none">• Log files from the <code>/mroot/etc/log/shelflog/</code> directory• Log files from the <code>/mroot/etc/log/acp/</code> directory• Event Management System (EMS) log data | The most recent lines of data up to a specified maximum. |

The content of AutoSupport messages can change between releases of ONTAP.

Files sent in weekly AutoSupport messages

Weekly AutoSupport messages contain additional configuration and status data that is useful to track changes in your system over time.

The following information is sent in weekly AutoSupport messages:

- Basic information about every subsystem
- Contents of selected `/mroot/etc` directory files
- Log files
- Output of commands that provide system information
- Additional information, including replicated database (RDB) information, service statistics, and more

How AutoSupport OnDemand obtains delivery instructions from technical support

AutoSupport OnDemand periodically communicates with technical support to obtain delivery instructions for sending, resending, and declining AutoSupport messages as well as uploading large files to the NetApp support site. AutoSupport OnDemand enables AutoSupport messages to be sent on-demand instead of waiting for the weekly AutoSupport job to run.

AutoSupport OnDemand consists of the following components:

- AutoSupport OnDemand client that runs on each node
- AutoSupport OnDemand service that resides in technical support

The AutoSupport OnDemand client periodically polls the AutoSupport OnDemand service to obtain delivery instructions from technical support. For example, technical support can use the AutoSupport OnDemand service to request that a new AutoSupport message be generated. When the AutoSupport OnDemand client polls the AutoSupport OnDemand service, the client obtains the delivery instructions and sends the new AutoSupport message on-demand as requested.

AutoSupport OnDemand is enabled by default. However, AutoSupport OnDemand relies on some AutoSupport settings to continue communicating with technical support. AutoSupport OnDemand automatically communicates with technical support when the following requirements are met:

- AutoSupport is enabled.
- AutoSupport is configured to send messages to technical support.
- AutoSupport is configured to use the HTTPS transport protocol.

The AutoSupport OnDemand client sends HTTPS requests to the same technical support location to which AutoSupport messages are sent. The AutoSupport OnDemand client does not accept incoming connections.



AutoSupport OnDemand uses the “autosupport” user account to communicate with technical support. ONTAP prevents you from deleting this account.

If you want to disable AutoSupport OnDemand, but keep AutoSupport enabled, use the command: `system node autosupport modify -ondemand-state disable`.

The following illustration shows how AutoSupport OnDemand sends HTTPS requests to technical support to obtain delivery instructions.



The delivery instructions can include requests for AutoSupport to do the following:

- Generate new AutoSupport messages.

Technical support might request new AutoSupport messages to help triage issues.

- Generate new AutoSupport messages that upload core dump files or performance archive files to the NetApp support site.

Technical support might request core dump or performance archive files to help triage issues.

- Retransmit previously generated AutoSupport messages.

This request automatically happens if a message was not received due to a delivery failure.

- Disable delivery of AutoSupport messages for specific trigger events.

Technical support might disable delivery of data that is not used.

Structure of AutoSupport messages sent by email

When an AutoSupport message is sent by email, the message has a standard subject, a brief body, and a large attachment in 7z file format that contains the data.



If AutoSupport is configured to hide private data, certain information, such as the hostname, is omitted or masked in the header, subject, body, and attachments.

Subject

The subject line of messages sent by the AutoSupport mechanism contains a text string that identifies the reason for the notification. The format of the subject line is as follows:

HA Group Notification from *System_Name* (*Message*) *Severity*

- *System_Name* is either the hostname or the system ID, depending on the AutoSupport configuration

Body

The body of the AutoSupport message contains the following information:

- Date and timestamp of the message
- Version of ONTAP on the node that generated the message

- System ID, serial number, and hostname of the node that generated the message
- AutoSupport sequence number
- SNMP contact name and location, if specified
- System ID and hostname of the HA partner node

Attached files

The key information in an AutoSupport message is contained in files that are compressed into a 7z file called `body.7z` and attached to the message.

The files contained in the attachment are specific to the type of AutoSupport message.

AutoSupport severity types

AutoSupport messages have severity types that help you understand the purpose of each message—for example, to draw immediate attention to an emergency problem, or only to provide information.

Messages have one of the following severities:

- **Alert:** Alert messages indicate that a next-higher level event might occur if you do not take some action.
You must take an action against alert messages within 24 hours.
- **Emergency:** Emergency messages are displayed when a disruption has occurred.
You must take an action against emergency messages immediately.
- **Error:** Error conditions indicate what might happen if you ignore.
- **Notice:** Normal but significant condition.
- **Info:** Informational message provides details about the issue, which you can ignore.
- **Debug:** Debug-level messages provide instructions you should perform.

If your internal support organization receives AutoSupport messages through email, the severity appears in the subject line of the email message.

Requirements for using AutoSupport

You should use HTTPS for delivery of AutoSupport messages to provide the best security and to support all of the latest AutoSupport features. Although AutoSupport supports HTTP and SMTP for delivery of AutoSupport messages, HTTPS is recommended.

Supported protocols

All of these protocols run on IPv4 or IPv6, based on the address family to which the name resolves.

| Protocol and port | Description |
|---------------------------------|--|
| HTTPS on port 443 | <p>This is the default protocol. You should use this whenever possible.</p> <p>This protocol supports AutoSupport OnDemand and uploads of large files.</p> <p>The certificate from the remote server is validated against the root certificate, unless you disable validation.</p> <p>The delivery uses an HTTP PUT request. With PUT, if the request fails during transmission, the request restarts where it stopped. If the server receiving the request does not support PUT, the delivery uses an HTTP POST request.</p> |
| HTTP on port 80 | <p>This protocol is preferred over SMTP.</p> <p>This protocol supports uploads of large files, but not AutoSupport OnDemand.</p> <p>The delivery uses an HTTP PUT request. With PUT, if the request fails during transmission, the request restarts where it stopped. If the server receiving the request does not support PUT, the delivery uses an HTTP POST request.</p> |
| SMTP on port 25 or another port | <p>You should use this protocol only if the network connection does not allow HTTPS or HTTP.</p> <p>The default port value is 25, but you can configure AutoSupport to use a different port.</p> <p>Keep the following limitations in mind when using SMTP:</p> <ul style="list-style-type: none"> • AutoSupport OnDemand and uploads of large files are not supported. • Data is not encrypted. <p>SMTP sends data in clear text, making text in the AutoSupport message easy to intercept and read.</p> <ul style="list-style-type: none"> • Limitations on message length and line length can be introduced. |

If you configure AutoSupport with specific email addresses for your internal support organization, or a support partner organization, those messages are always sent by SMTP.

For example, if you use the recommended protocol to send messages to technical support and you also want to send messages to your internal support organization, your messages will be transported using both HTTPS

and SMTP, respectively.

AutoSupport limits the maximum file size for each protocol. The default setting for HTTP and HTTPS transfers is 25 MB. The default setting for SMTP transfers is 5 MB. If the size of the AutoSupport message exceeds the configured limit, AutoSupport delivers as much of the message as possible. You can edit the maximum size by modifying AutoSupport configuration. See the `system node autosupport modify` man page for more information.



AutoSupport automatically overrides the maximum file size limit for the HTTPS and HTTP protocols when you generate and send AutoSupport messages that upload core dump or performance archive files to the NetApp support site or a specified URI. The automatic override applies only when you upload files by using the `system node autosupport invoke-core-upload` or the `system node autosupport invoke-performance-archive` commands.

Configuration requirements

Depending on your network configuration, use of HTTP or HTTPS protocols may require additional configuration of a proxy URL. If you use HTTP or HTTPS to send AutoSupport messages to technical support and you have a proxy, you must identify the URL for that proxy. If the proxy uses a port other than the default port, which is 3128, you can specify the port for that proxy. You can also specify a user name and password for proxy authentication.

If you use SMTP to send AutoSupport messages either to your internal support organization or to technical support, you must configure an external mail server. The storage system does not function as a mail server; it requires an external mail server at your site to send mail. The mail server must be a host that listens on the SMTP port (25) or another port, and it must be configured to send and receive 8-bit Multipurpose Internet Mail Extensions (MIME) encoding. Example mail hosts include a UNIX host running an SMTP server such as the sendmail program and a Windows server running the Microsoft Exchange server. You can have one or more mail hosts.

Set up AutoSupport

You can control whether and how AutoSupport information is sent to technical support and your internal support organization, and then test that the configuration is correct.

About this task

In ONTAP 9.5 and later releases, you can enable AutoSupport and modify its configuration on all nodes of the cluster simultaneously. When a new node joins the cluster, the node inherits the AutoSupport cluster configuration automatically. You do not have to update the configuration on each node separately.



Beginning with ONTAP 9.5, the scope of the `system node autosupport modify` command is cluster-wide. The AutoSupport configuration is modified on all nodes in the cluster, even when the `-node` option is specified. The option is ignored, but it has been retained for CLI backward compatibility.

In ONTAP 9.4 and earlier releases, the scope of the "system node autosupport modify" command is specific to the node. The AutoSupport configuration should be modified on each node in your cluster.

By default, AutoSupport is enabled on each node to send messages to technical support by using the HTTPS transport protocol.

Steps

1. Ensure that AutoSupport is enabled:

```
system node autosupport modify -state enable
```

2. If you want technical support to receive AutoSupport messages, use the following command:

```
system node autosupport modify -support enable
```

You must enable this option if you want to enable AutoSupport to work with AutoSupport OnDemand or if you want to upload large files, such as core dump and performance archive files, to technical support or a specified URL.

3. If technical support is enabled to receive AutoSupport messages, specify which transport protocol to use for the messages.

You can choose from the following options:

| If you want to... | Then set the following parameters of the <code>system node autosupport modify</code> command... |
|--------------------------------------|--|
| Use the default HTTPS protocol | <ol style="list-style-type: none">a. Set <code>-transport</code> to <code>https</code>.b. If you use a proxy, set <code>-proxy-url</code> to the URL of your proxy. This configuration supports communication with AutoSupport OnDemand and uploads of large files. |
| Use HTTP that is preferred over SMTP | <ol style="list-style-type: none">a. Set <code>-transport</code> to <code>http</code>.b. If you use a proxy, set <code>-proxy-url</code> to the URL of your proxy. This configuration supports uploads of large files, but not AutoSupport OnDemand. |
| Use SMTP | <p>Set <code>-transport</code> to <code>smtp</code>.</p> <p>This configuration does not support AutoSupport OnDemand or uploads of large files.</p> |

4. If you want your internal support organization or a support partner to receive AutoSupport messages, perform the following actions:

- a. Identify the recipients in your organization by setting the following parameters of the `system node autosupport modify` command:

| Set this parameter... | To this... |
|-----------------------|------------|
|-----------------------|------------|

| | |
|-------------------------------|---|
| <code>-to</code> | Up to five comma-separated individual email addresses or distribution lists in your internal support organization that will receive key AutoSupport messages |
| <code>-noteto</code> | Up to five comma-separated individual email addresses or distribution lists in your internal support organization that will receive a shortened version of key AutoSupport messages designed for cell phones and other mobile devices |
| <code>-partner-address</code> | Up to five comma-separated individual email addresses or distribution lists in your support partner organization that will receive all AutoSupport messages |

- b. Check that addresses are correctly configured by listing the destinations using the `system node autosupport destinations show` command.
5. If you are sending messages to your internal support organization or you chose SMTP transport for messages to technical support, configure SMTP by setting the following parameters of the `system node autosupport modify` command:

- Set `-mail-hosts` to one or more mail hosts, separated by commas.

You can set a maximum of five.

You can configure a port value for each mail host by specifying a colon and port number after the mail host name: for example, `mymailhost.example.com:5678`, where 5678 is the port for the mail host.

- Set `-from` to the email address that sends the AutoSupport message.

6. Configure DNS.
7. (Optional) Add command options if you want to change specific settings:

| If you want to do this... | Then set the following parameters of the <code>system node autosupport modify</code> command... |
|--|---|
| Hide private data by removing, masking, or encoding sensitive data in the messages | Set <code>-remove-private-data</code> to <code>true</code> . If you change from <code>false</code> to <code>true</code> , all AutoSupport history and all associated files are deleted. |
| Stop sending performance data in periodic AutoSupport messages | Set <code>-perf</code> to <code>false</code> . |

8. Check the overall configuration by using the `system node autosupport show` command with the `-node` parameter.
9. Verify the AutoSupport operation by using the `system node autosupport check show` command.

If any problems are reported, use the `system node autosupport check show-details` command to view more information.

10. Test that AutoSupport messages are being sent and received:

- a. Use the `system node autosupport invoke` command with the `-type` parameter set to `test`.

```
cluster1::> system node autosupport invoke -type test -node node1
```

- b. Confirm that NetApp is receiving your AutoSupport messages:

```
system node autosupport history show -node local
```

The status of the latest outgoing AutoSupport message should eventually change to `sent-successful` for all appropriate protocol destinations.

- c. (Optional) Confirm that the AutoSupport message is being sent to your internal support organization or to your support partner by checking the email of any address that you configured for the `-to`, `-noteto`, or `-partner-address` parameters of the `system node autosupport modify` command.

Upload core dump files

When a core dump file is saved, an event message is generated. If the AutoSupport service is enabled and configured to send messages to NetApp support, an AutoSupport message is transmitted, and an automated email acknowledgement is sent to you.

What you'll need

- You must have set up AutoSupport with the following settings:
 - AutoSupport is enabled on the node.
 - AutoSupport is configured to send messages to technical support.
 - AutoSupport is configured to use the HTTP or HTTPS transport protocol.

The SMTP transport protocol is not supported when sending messages that include large files, such as core dump files.

About this task

You can also upload the core dump file through the AutoSupport service over HTTPS by using the `system node autosupport invoke-core-upload` command, if requested by NetApp support.

[How to upload a file to NetApp](#)

Steps

1. View the core dump files for a node by using the `system node coredump show` command.

In the following example, core dump files are displayed for the local node:

```
cluster1::> system node coredump show -node local
Node:Type Core Name Saved Panic Time
-----
node:kernel
core.4073000068.2013-09-11.15_05_01.nz true 9/11/2013 15:05:01
```

2. Generate an AutoSupport message and upload a core dump file by using the `system node autosupport invoke-core-upload` command.

In the following example, an AutoSupport message is generated and sent to the default location, which is technical support, and the core dump file is uploaded to the default location, which is the NetApp support site:

```
cluster1::> system node autosupport invoke-core-upload -core-filename
core.4073000068.2013-09-11.15_05_01.nz -node local
```

In the following example, an AutoSupport message is generated and sent to the location specified in the URI, and the core dump file is uploaded to the URI:

```
cluster1::> system node autosupport invoke-core-upload -uri
https://files.company.com -core-filename
core.4073000068.2013-09-11.15_05_01.nz -node local
```

Upload performance archive files

You can generate and send an AutoSupport message that contains a performance archive. By default, NetApp technical support receives the AutoSupport message, and the performance archive is uploaded to the NetApp support site. You can specify an alternate destination for the message and upload.

What you'll need

- You must have set up AutoSupport with the following settings:
 - AutoSupport is enabled on the node.
 - AutoSupport is configured to send messages to technical support.
 - AutoSupport is configured to use the HTTP or HTTPS transport protocol.

The SMTP transport protocol is not supported when sending messages that include large files, such as performance archive files.

About this task

You must specify a start date for the performance archive data that you want to upload. Most storage systems retain performance archives for two weeks, enabling you to specify a start date up to two weeks ago. For example, if today is January 15, you can specify a start date of January 2.

Step

1. Generate an AutoSupport message and upload the performance archive file by using the `system node autosupport invoke-performance-archive` command.

In the following example, 4 hours of performance archive files from January 12, 2015 are added to an AutoSupport message and uploaded to the default location, which is the NetApp support site:

```
cluster1::> system node autosupport invoke-performance-archive -node
local -start-date 1/12/2015 13:42:09 -duration 4h
```

In the following example, 4 hours of performance archive files from January 12, 2015 are added to an AutoSupport message and uploaded to the location specified by the URI:

```
cluster1::> system node autosupport invoke-performance-archive -node
local -start-date 1/12/2015 13:42:09 -duration 4h -uri
https://files.company.com
```

Get AutoSupport message descriptions

The descriptions of the AutoSupport messages that you receive are available through the ONTAP Syslog Translator.

Steps

1. Go to the [Syslog Translator](#).
2. In the **Release** field, enter the version of ONTAP you are using. In the **Search String** field, enter "callhome". Select **Translate**.
3. The Syslog Translator will alphabetically list all events that match the message string you entered.

Commands for managing AutoSupport

You use the `system node autosupport` commands to change or view AutoSupport configuration, display information about previous AutoSupport messages, and send, resend or cancel an AutoSupport message.

Configure AutoSupport

| If you want to... | Use this command... |
|--|--|
| Control whether any AutoSupport messages are sent | <code>system node autosupport modify</code> with the <code>-state</code> parameter |
| Control whether AutoSupport messages are sent to technical support | <code>system node autosupport modify</code> with the <code>-support</code> parameter |

| If you want to... | Use this command... |
|--|---|
| Set up AutoSupport or modify the configuration of AutoSupport | <code>system node autosupport modify</code> |
| Enable and disable AutoSupport messages to your internal support organization for individual trigger events, and specify additional subsystem reports to include in messages sent in response to individual trigger events | <code>system node autosupport trigger modify</code> |

Display information about the AutoSupport configuration

| If you want to... | Use this command... |
|---|---|
| Display the AutoSupport configuration | <code>system node autosupport show</code> with the <code>-node</code> parameter |
| View a summary of all addresses and URLs that receive AutoSupport messages | <code>system node autosupport destinations show</code> |
| Display which AutoSupport messages are sent to your internal support organization for individual trigger events | <code>system node autosupport trigger show</code> |
| Display status of AutoSupport configuration as well as delivery to various destinations | <code>system node autosupport check show</code> |
| Display detailed status of AutoSupport configuration as well as delivery to various destinations | <code>system node autosupport check show-details</code> |

Display information about past AutoSupport messages

| If you want to... | Use this command... |
|---|--|
| Display information about one or more of the 50 most recent AutoSupport messages | <code>system node autosupport history show</code> |
| Display information about recent AutoSupport messages generated to upload core dump or performance archive files to the technical support site or a specified URI | <code>system node autosupport history show-upload-details</code> |
| View the information in the AutoSupport messages including the name and size of each file collected for the message along with any errors | <code>system node autosupport manifest show</code> |

Send, resend, or cancel AutoSupport messages

| If you want to... | Use this command... |
|--|---|
| <p>Retransmit a locally stored AutoSupport message, identified by its AutoSupport sequence number</p> <div><p>If you retransmit an AutoSupport message, and if support already received that message, the support system will not create a duplicate case. If, on the other hand, support did not receive that message, then the AutoSupport system will analyze the message and create a case, if necessary.</p></div> | <pre>system node autosupport history retransmit</pre> |
| <p>Generate and send an AutoSupport message—for example, for testing purposes</p> | <pre>system node autosupport invoke</pre> <div><p>Use the <code>-force</code> parameter to send a message even if AutoSupport is disabled. Use the <code>-uri</code> parameter to send the message to the destination you specify instead of the configured destination.</p></div> |
| <p>Cancel an AutoSupport message</p> | <pre>system node autosupport history cancel</pre> |

Related information

[ONTAP 9 Commands](#)

Information included in the AutoSupport manifest

The AutoSupport manifest provides you with a detailed view of the files collected for each AutoSupport message. The AutoSupport manifest also includes information about collection errors when AutoSupport cannot collect the files it needs.

The AutoSupport manifest includes the following information:

- Sequence number of the AutoSupport message
- Which files AutoSupport included in the AutoSupport message
- Size of each file, in bytes
- Status of the AutoSupport manifest collection
- Error description, if AutoSupport failed to collect one or more files

You can view the AutoSupport manifest by using the `system node autosupport manifest show` command.

The AutoSupport manifest is included with every AutoSupport message and presented in XML format, which

means that you can either use a generic XML viewer to read it or view it using the Active IQ (formerly known as My AutoSupport) portal.

AutoSupport case suppression during scheduled maintenance windows

AutoSupport case suppression enables you to stop unnecessary cases from being created by AutoSupport messages that are triggered during scheduled maintenance windows.

To suppress AutoSupport cases, you must manually invoke an AutoSupport message with a specially formatted text string: `MAINT=xh`. `x` is the duration of the maintenance window in units of hours.

Related information

[How to suppress automatic case creation during scheduled maintenance windows](#)

Troubleshoot AutoSupport

Troubleshoot AutoSupport when messages are not received

If the system does not send the AutoSupport message, you can determine whether that is because AutoSupport cannot generate the message or cannot deliver the message.

Steps

1. Check delivery status of the messages by using the `system node autosupport history show` command.
2. Read the status.

| This status | Means |
|------------------------|---|
| initializing | The collection process is starting. If this state is temporary, all is well. However, if this state persists, there is an issue. |
| collection-failed | AutoSupport cannot create the AutoSupport content in the spool directory. You can view what AutoSupport is trying to collect by entering the <code>system node autosupport history show -detail</code> command. |
| collection-in-progress | AutoSupport is collecting AutoSupport content. You can view what AutoSupport is collecting by entering the <code>system node autosupport manifest show</code> command. |
| queued | AutoSupport messages are queued for delivery, but not yet delivered. |
| transmitting | AutoSupport is currently delivering messages. |
| sent-successful | AutoSupport successfully delivered the message. You can find out where AutoSupport delivered the message by entering the <code>system node autosupport history show -delivery</code> command. |

| This status | Means |
|---------------------|---|
| ignore | AutoSupport has no destinations for the message. You can view the delivery details by entering the <code>system node autosupport history show -delivery</code> command. |
| re-queued | AutoSupport tried to deliver messages, but the attempt failed. As a result, AutoSupport placed the messages back in the delivery queue for another attempt. You can view the error by entering the <code>system node autosupport history show</code> command. |
| transmission-failed | AutoSupport failed to deliver the message the specified number of times and stopped trying to deliver the message. You can view the error by entering the <code>system node autosupport history show</code> command. |
| ondemand-ignore | The AutoSupport message was processed successfully, but the AutoSupport OnDemand service chose to ignore it. |

3. Perform one of the following actions:

| For this status | Do this |
|---|---|
| initializing or collection-failed | <p>Contact NetApp Support, because AutoSupport cannot generate the message. Mention the following Knowledge Base article:</p> <p>AutoSupport is failing to deliver: status is stuck in initializing</p> |
| ignore, re-queued, or transmission failed | Check that destinations are correctly configured for SMTP, HTTP, or HTTPS because AutoSupport cannot deliver the message. |

Troubleshoot AutoSupport message delivery over HTTP or HTTPS

If the system does not send the expected AutoSupport message and you are using HTTP or HTTPS, or the Automatic Update feature is not working, you can check a number of settings to resolve the problem.

What you'll need

You should have confirmed basic network connectivity and DNS lookup:

- Your node management LIF must be up for operational and administrative status.
- You must be able to ping a functioning host on the same subnet from the cluster management LIF (not a LIF on any of the nodes).
- You must be able to ping a functioning host outside the subnet from the cluster management LIF.
- You must be able to ping a functioning host outside the subnet from the cluster management LIF using the name of the host (not the IP address).

About this task

These steps are for cases when you have determined that AutoSupport can generate the message, but cannot deliver the message over HTTP or HTTPS.

If you encounter errors or cannot complete a step in this procedure, determine and address the problem before proceeding to the next step.

Steps

1. Display the detailed status of the AutoSupport subsystem:

```
system node autosupport check show-details
```

This includes verifying connectivity to AutoSupport destinations by sending test messages and providing a list of possible errors in your AutoSupport configuration settings.

2. Verify the status of the node management LIF:

```
network interface show -home-node local -role node-mgmt -fields  
vserver,lif,status-oper,status-admin,address,role
```

The `status-oper` and `status-admin` fields should return “up”.

3. Record the SVM name, the LIF name, and the LIF IP address for later use.
4. Ensure that DNS is enabled and configured correctly:

```
vserver services name-service dns show
```

5. Address any errors returned by the AutoSupport message:

```
system node autosupport history show -node * -fields node,seq-  
num,destination,last-update,status,error
```

For assistance troubleshooting any returned errors, see the [ONTAP AutoSupport \(Transport HTTPS and HTTP\) Resolution Guide](#).

6. Confirm that the cluster can access both the servers it needs and the Internet successfully:

- a. `network traceroute -lif node-management_LIF -destination DNS server`
- b. `network traceroute -lif node_management_LIF -destination support.netapp.com`



The address `support.netapp.com` itself does not respond to ping/traceroute, but the per-hop information is valuable.

- c. `system node autosupport show -fields proxy-url`
- d. `network traceroute -node node_management_LIF -destination proxy_url`

If any of these routes are not functioning, try the same route from a functioning host on the same subnet as the cluster, using the “traceroute” or “tracert” utility found on most third-party network clients. This assists you in determining whether the issue is in your network configuration or your cluster configuration.

7. If you are using HTTPS for your AutoSupport transport protocol, ensure that HTTPS traffic can exit your network:

- a. Configure a web client on the same subnet as the cluster management LIF.

Ensure that all configuration parameters are the same values as for the AutoSupport configuration, including using the same proxy server, user name, password, and port.

- b. Access `https://support.netapp.com` with the web client.

The access should be successful. If not, ensure that all firewalls are configured correctly to allow HTTPS and DNS traffic, and that the proxy server is configured correctly. For more information on configuring static name resolution for `support.netapp.com`, see the Knowledge Base article [How would a HOST entry be added in ONTAP for support.netapp.com?](#)

8. Beginning with ONTAP 9.10.1, if you enabled the Automatic Update feature, ensure you have HTTPS connectivity to the following additional URLs:

- `https://support-sg-emea.netapp.com`
- `https://support-sg-naeast.netapp.com`
- `https://support-sg-nawest.netapp.com`

Troubleshoot AutoSupport message delivery over SMTP

If the system cannot deliver AutoSupport messages over SMTP, you can check a number of settings to resolve the problem.

What you'll need

You should have confirmed basic network connectivity and DNS lookup:

- Your node management LIF must be up for operational and administrative status.
- You must be able to ping a functioning host on the same subnet from the cluster management LIF (not a LIF on any of the nodes).
- You must be able to ping a functioning host outside the subnet from the cluster management LIF.
- You must be able to ping a functioning host outside the subnet from the cluster management LIF using the name of the host (not the IP address).

About this task

These steps are for cases when you have determined that AutoSupport can generate the message, but cannot deliver the message over SMTP.

If you encounter errors or cannot complete a step in this procedure, determine and address the problem before proceeding to the next step.

All commands are entered at the ONTAP command-line interface, unless otherwise specified.

Steps

1. Verify the status of the node management LIF:

```
network interface show -home-node local -role node-mgmt -fields
vservers,lif,status-oper,status-admin,address,role
```

The `status-oper` and `status-admin` fields should return `up`.

2. Record the SVM name, the LIF name, and the LIF IP address for later use.
3. Ensure that DNS is enabled and configured correctly:

```
vserver services name-service dns show
```

4. Display all of the servers configured to be used by AutoSupport:

```
system node autosupport show -fields mail-hosts
```

Record all server names displayed.

5. For each server displayed by the previous step, and `support.netapp.com`, ensure that the server or URL can be reached by the node:

```
network traceroute -node local -destination server_name
```

If any of these routes is not functioning, try the same route from a functioning host on the same subnet as the cluster, using the “traceroute” or “tracert” utility found on most third-party network clients. This assists you in determining whether the issue is in your network configuration or your cluster configuration.

6. Log in to the host designated as the mail host, and ensure that it can serve SMTP requests:

```
netstat -aAn|grep 25
```

25 is the listener SMTP port number.

A message similar to the following text is displayed:

```
ff64878c tcp          0      0 *.25    *.*    LISTEN.
```

7. From some other host, open a Telnet session with the SMTP port of the mail host:

```
telnet mailhost 25
```

A message similar to the following text is displayed:

```
220 filer.yourco.com Sendmail 4.1/SMI-4.1 ready at Thu, 30 Nov 2014
10:49:04 PST
```

8. At the telnet prompt, ensure that a message can be relayed from your mail host:

```
HELO domain_name
```

```
MAIL FROM: your_email_address
```

```
RCPT TO: autosupport@netapp.com
```

`domain_name` is the domain name of your network.

If an error is returned saying that relaying is denied, relaying is not enabled on the mail host. Contact your

system administrator.

9. At the telnet prompt, send a test message:

DATA

SUBJECT: TESTING THIS IS A TEST

.



Ensure that you enter the last period (.) on a line by itself. The period indicates to the mail host that the message is complete.

If an error is returned, your mail host is not configured correctly. Contact your system administrator.

10. From the ONTAP command-line interface, send an AutoSupport test message to a trusted email address that you have access to:

```
system node autosupport invoke -node local -type test
```

11. Find the sequence number of the attempt:

```
system node autosupport history show -node local -destination smtp
```

Find the sequence number for your attempt based on the timestamp. It is probably the most recent attempt.

12. Display the error for your test message attempt:

```
system node autosupport history show -node local -seq-num seq_num -fields error
```

If the error displayed is `Login denied`, your SMTP server is not accepting send requests from the cluster management LIF. If you do not want to change to using HTTPS as your transport protocol, contact your site network administrator to configure the SMTP gateways to address this issue.

If this test succeeds but the same message sent to `mailto:autosupport@netapp.com` does not, ensure that SMTP relay is enabled on all of your SMTP mail hosts, or use HTTPS as a transport protocol.

If even the message to the locally administered email account does not succeed, confirm that your SMTP servers are configured to forward attachments with both of these characteristics:

- The “7z” suffix
- The “application/x-7x-compressed” MIME type.

Troubleshoot the AutoSupport subsystem

The `system node check show` commands can be used to verify and troubleshoot any issues related to the AutoSupport configuration and delivery.

Step

1. Use the following commands to display the status of the AutoSupport subsystem.

| Use this command... | To do this... |
|---|---|
| <code>system node autosupport check show</code> | Display overall status of the AutoSupport subsystem, such as the status of AutoSupport HTTP or HTTPS destination, AutoSupport SMTP destinations, AutoSupport OnDemand Server, and AutoSupport configuration |
| <code>system node autosupport check show-details</code> | Display detailed status of the AutoSupport subsystem, such as detailed descriptions of errors and the corrective actions |

Monitor the health of your system

Monitor the health of your system overview

Health monitors proactively monitor certain critical conditions in your cluster and raise alerts if they detect a fault or risk. If there are active alerts, the system health status reports a degraded status for the cluster. The alerts include the information that you need to respond to degraded system health.

If the status is degraded, you can view details about the problem, including the probable cause and recommended recovery actions. After you resolve the problem, the system health status automatically returns to OK.

The system health status reflects multiple separate health monitors. A degraded status in an individual health monitor causes a degraded status for the overall system health.

For details on how ONTAP supports cluster switches for system health monitoring in your cluster, you can refer to the *Hardware Universe*.

[Supported switches in the Hardware Universe](#)

For details on the causes of Cluster Switch Health Monitor (CSHM) AutoSupport messages, and the necessary actions required to resolve these alerts, you can refer to the Knowledgebase article.

[AutoSupport Message: Health Monitor Process CSHM](#)

How health monitoring works

Individual health monitors have a set of policies that trigger alerts when certain conditions occur. Understanding how health monitoring works can help you respond to problems and control future alerts.

Health monitoring consists of the following components:

- Individual health monitors for specific subsystems, each of which has its own health status

For example, the Storage subsystem has a node connectivity health monitor.

- An overall system health monitor that consolidates the health status of the individual health monitors

A degraded status in any single subsystem results in a degraded status for the entire system. If no subsystems have alerts, the overall system status is OK.

Each health monitor is made up of the following key elements:

- Alerts that the health monitor can potentially raise

Each alert has a definition, which includes details such as the severity of the alert and its probable cause.

- Health policies that identify when each alert is triggered

Each health policy has a rule expression, which is the exact condition or change that triggers the alert.

A health monitor continuously monitors and validates the resources in its subsystem for condition or state changes. When a condition or state change matches a rule expression in a health policy, the health monitor raises an alert. An alert causes the subsystem's health status and the overall system health status to become degraded.

Ways to respond to system health alerts

When a system health alert occurs, you can acknowledge it, learn more about it, repair the underlying condition, and prevent it from occurring again.

When a health monitor raises an alert, you can respond in any of the following ways:

- Get information about the alert, which includes the affected resource, alert severity, probable cause, possible effect, and corrective actions.
- Get detailed information about the alert, such as the time when the alert was raised and whether anyone else has acknowledged the alert already.
- Get health-related information about the state of the affected resource or subsystem, such as a specific shelf or disk.
- Acknowledge the alert to indicate that someone is working on the problem, and identify yourself as the "Acknowledger."
- Resolve the problem by taking the corrective actions provided in the alert, such as fixing cabling to resolve a connectivity problem.
- Delete the alert, if the system did not automatically clear it.
- Suppress an alert to prevent it from affecting the health status of a subsystem.

Suppressing is useful when you understand a problem. After you suppress an alert, it can still occur, but the subsystem health displays as "ok-with-suppressed." when the suppressed alert occurs.

System health alert customization

You can control which alerts a health monitor generates by enabling and disabling the system health policies that define when alerts are triggered. This enables you to customize the health monitoring system for your particular environment.

You can learn the name of a policy either by displaying detailed information about a generated alert or by displaying policy definitions for a specific health monitor, node, or alert ID.

Disabling health policies is different from suppressing alerts. When you suppress an alert, it does not affect the subsystem's health status, but the alert can still occur.

If you disable a policy, the condition or state that is defined in its policy rule expression no longer triggers an alert.

Example of an alert that you want to disable

For example, suppose an alert occurs that is not useful to you. You use the `system health alert show -instance` command to obtain the Policy ID for the alert. You use the policy ID in the `system health policy definition show` command to view information about the policy. After reviewing the rule expression and other information about the policy, you decide to disable the policy. You use the `system health policy definition modify` command to disable the policy.

How health alerts trigger AutoSupport messages and events

System health alerts trigger AutoSupport messages and events in the Event Management System (EMS), enabling you to monitor the health of the system using AutoSupport messages and the EMS in addition to using the health monitoring system directly.

Your system sends an AutoSupport message within five minutes of an alert. The AutoSupport message includes all alerts generated since the previous AutoSupport message, except for alerts that duplicate an alert for the same resource and probable cause within the previous week.

Some alerts do not trigger AutoSupport messages. An alert does not trigger an AutoSupport message if its health policy disables the sending of AutoSupport messages. For example, a health policy might disable AutoSupport messages by default because AutoSupport already generates a message when the problem occurs. You can configure policies to not trigger AutoSupport messages by using the `system health policy definition modify` command.

You can view a list of all of the alert-triggered AutoSupport messages sent in the previous week using the `system health autosupport trigger history show` command.

Alerts also trigger the generation of events to the EMS. An event is generated each time an alert is created and each time an alert is cleared.

Available cluster health monitors

There are several health monitors that monitor different parts of a cluster. Health monitors help you to recover from errors within ONTAP systems by detecting events, sending alerts to you, and deleting events as they clear.

| Health monitor name (identifier) | Subsystem name (identifier) | Purpose |
|--------------------------------------|--|---|
| Cluster switch(cluster-switch) | Switch (Switch-Health) | <p>Monitors cluster network switches and management network switches for temperature, utilization, interface configuration, redundancy (cluster network switches only), and fan and power supply operation. The cluster switch health monitor communicates with switches through SNMP. SNMPv2c is the default setting.</p> <div>  <p>Beginning with ONTAP 9.2, this monitor can detect and report when a cluster switch has rebooted since the last polling period.</p> </div> |
| MetroCluster Fabric | Switch | Monitors the MetroCluster configuration back-end fabric topology and detects misconfigurations such as incorrect cabling and zoning, and ISL failures. |
| MetroCluster Health | Interconnect, RAID, and storage | Monitors FC-VI adapters, FC initiator adapters, left-behind aggregates and disks, and inter-cluster ports |
| Node connectivity(node-connect) | CIFS nondisruptive operations (CIFS-NDO) | Monitors SMB connections for nondisruptive operations to Hyper-V applications. |
| | Storage (SAS-connect) | Monitors shelves, disks, and adapters at the node level for appropriate paths and connections. |
| System | not applicable | Aggregates information from other health monitors. |
| System connectivity (system-connect) | Storage (SAS-connect) | Monitors shelves at the cluster level for appropriate paths to two HA clustered nodes. |

Receive system health alerts automatically

You can manually view system health alerts by using the `system health alert show` command. However, you should subscribe to specific Event Management System (EMS) messages to automatically receive notifications when a health monitor generates an alert.

About this task

The following procedure shows you how to set up notifications for all `hm.alert.raised` messages and all `hm.alert.cleared` messages.

All `hm.alert.raised` messages and all `hm.alert.cleared` messages include an SNMP trap. The names of the SNMP traps are `HealthMonitorAlertRaised` and `HealthMonitorAlertCleared`. For information about SNMP traps, see the *Network Management Guide*.

Steps

1. Use the `event destination create` command to define the destination to which you want to send the EMS messages.

```
cluster1::> event destination create -name health_alerts -mail  
admin@example.com
```

2. Use the `event route add-destinations` command to route the `hm.alert.raised` message and the `hm.alert.cleared` message to a destination.

```
cluster1::> event route add-destinations -messagename hm.alert*  
-destinations health_alerts
```

Related information

[Network management](#)

Respond to degraded system health

When your system's health status is degraded, you can show alerts, read about the probable cause and corrective actions, show information about the degraded subsystem, and resolve the problem. Suppressed alerts are also shown so that you can modify them and see whether they have been acknowledged.

About this task

You can discover that an alert was generated by viewing an AutoSupport message or an EMS event, or by using the `system health` commands.

Steps

1. Use the `system health alert show` command to view the alerts that are compromising the system's health.
2. Read the alert's probable cause, possible effect, and corrective actions to determine whether you can resolve the problem or need more information.

3. If you need more information, use the `system health alert show -instance` command to view additional information available for the alert.
4. Use the `system health alert modify` command with the `-acknowledge` parameter to indicate that you are working on a specific alert.
5. Take corrective action to resolve the problem as described by the `Corrective Actions` field in the alert.

The corrective actions might include rebooting the system.

When the problem is resolved, the alert is automatically cleared. If the subsystem has no other alerts, the health of the subsystem changes to `OK`. If the health of all subsystems is `OK`, the overall system health status changes to `OK`.

6. Use the `system health status show` command to confirm that the system health status is `OK`.

If the system health status is not `OK`, repeat this procedure.

Example of responding to degraded system health

By reviewing a specific example of degraded system health caused by a shelf that lacks two paths to a node, you can see what the CLI displays when you respond to an alert.

After starting ONTAP, you check the system health and you discover that the status is degraded:

```
cluster1::>system health status show
Status
-----
degraded
```

You show alerts to find out where the problem is, and see that shelf 2 does not have two paths to node1:

```
cluster1::>system health alert show
      Node: node1
      Resource: Shelf ID 2
      Severity: Major
      Indication Time: Mon Nov 10 16:48:12 2013
      Probable Cause: Disk shelf 2 does not have two paths to controller
                      node1.
      Possible Effect: Access to disk shelf 2 via controller node1 will be
                      lost with a single hardware component failure (e.g.
                      cable, HBA, or IOM failure).
      Corrective Actions: 1. Halt controller node1 and all controllers attached
to disk shelf 2.
                        2. Connect disk shelf 2 to controller node1 via two
paths following the rules in the Universal SAS and ACP Cabling Guide.
                        3. Reboot the halted controllers.
                        4. Contact support personnel if the alert persists.
```

You display details about the alert to get more information, including the alert ID:

```

cluster1::>system health alert show -monitor node-connect -alert-id
DualPathToDiskShelf_Alert -instance
    Node: node1
    Monitor: node-connect
    Alert ID: DualPathToDiskShelf_Alert
    Alerting Resource: 50:05:0c:c1:02:00:0f:02
    Subsystem: SAS-connect
    Indication Time: Mon Mar 21 10:26:38 2011
    Perceived Severity: Major
    Probable Cause: Connection_establishment_error
    Description: Disk shelf 2 does not have two paths to controller
node1.
    Corrective Actions: 1. Halt controller node1 and all controllers
attached to disk shelf 2.
                        2. Connect disk shelf 2 to controller node1 via
two paths following the rules in the Universal SAS and ACP Cabling Guide.
                        3. Reboot the halted controllers.
                        4. Contact support personnel if the alert
persists.
    Possible Effect: Access to disk shelf 2 via controller node1 will
be lost with a single
    hardware component failure (e.g. cable, HBA, or IOM failure).
    Acknowledge: false
    Suppress: false
    Policy: DualPathToDiskShelf_Policy
    Acknowledger: -
    Suppressor: -
    Additional Information: Shelf uuid: 50:05:0c:c1:02:00:0f:02
                        Shelf id: 2
                        Shelf Name: 4d.shelf2
                        Number of Paths: 1
                        Number of Disks: 6
                        Adapter connected to IOMA:
                        Adapter connected to IOMB: 4d
    Alerting Resource Name: Shelf ID 2

```

You acknowledge the alert to indicate that you are working on it.

```

cluster1::>system health alert modify -node node1 -alert-id
DualPathToDiskShelf_Alert -acknowledge true

```

You fix the cabling between shelf 2 and node1, and then reboot the system. Then you check system health again, and see that the status is OK:

```
cluster1::>system health status show
Status
-----
OK
```

Configure discovery of cluster and management network switches

The cluster switch health monitor automatically attempts to discover your cluster and management network switches using the Cisco Discovery Protocol (CDP). You must configure the health monitor if it cannot automatically discover a switch or if you do not want to use CDP for automatic discovery.

About this task

The `system cluster-switch show` command lists the switches that the health monitor discovered. If you do not see a switch that you expected to see in that list, then the health monitor cannot automatically discover it.

Steps

1. If you want to use CDP for automatic discovery, do the following:
 - a. Ensure that the Cisco Discovery Protocol (CDP) is enabled on your switches.

Refer to your switch documentation for instructions.
 - b. Run the following command on each node in the cluster to verify whether CDP is enabled or disabled:

```
run -node node_name -command options cdpd.enable
```

If CDP is enabled, go to step d. If CDP is disabled, go to step c.

- c. Run the following command to enable CDP:

```
run -node node_name -command options cdpd.enable on
```

Wait five minutes before you go to the next step.

- d. Use the `system cluster-switch show` command to verify whether ONTAP can now automatically discover the switches.
2. If the health monitor cannot automatically discover a switch, use the `system cluster-switch create` command to configure discovery of the switch:

```
cluster1::> system cluster-switch create -device switch1 -address
192.0.2.250 -snmp-version SNMPv2c -community cshml! -model NX5020 -type
cluster-network
```

Wait five minutes before you go to the next step.

3. Use the `system cluster-switch show` command to verify that ONTAP can discover the switch for

which you added information.

After you finish

Verify that the health monitor can monitor your switches.

Verify the monitoring of cluster and management network switches

The cluster switch health monitor automatically attempts to monitor the switches that it discovers; however, monitoring might not happen automatically if the switches are not configured correctly. You should verify that the health monitor is properly configured to monitor your switches.

Steps

1. To identify the switches that the cluster switch health monitor discovered, enter the following command:

ONTAP 9.8 and later

```
system switch ethernet show
```

ONTAP 9.7 and earlier

```
system cluster-switch show
```

If the `Model` column displays the value `OTHER`, then ONTAP cannot monitor the switch. ONTAP sets the value to `OTHER` if a switch that it automatically discovers is not supported for health monitoring.



If a switch does not display in the command output, you must configure discovery of the switch.

2. Upgrade to the latest supported switch software and reference the configuration file (RCF) from the NetApp Support Site.

[NetApp Support Downloads page](#)

The community string in the switch's RCF must match the community string that the health monitor is configured to use. By default, the health monitor uses the community string `cshml!`.



At this time, the health monitor only supports SNMPv2.

If you need to change information about a switch that the cluster monitors, you can modify the community string that the health monitor uses by using the following command:

ONTAP 9.8 and later

```
system switch ethernet modify
```

ONTAP 9.7 and earlier

```
system cluster-switch modify
```

3. Verify that the switch's management port is connected to the management network.

This connection is required to perform SNMP queries.

Commands for monitoring the health of your system

You can use the `system health` commands to display information about the health of system resources, to respond to alerts, and to configure future alerts. Using the CLI commands enables you to view in-depth information about how health monitoring is configured. The man pages for the commands contain more information.

Display the status of system health

| If you want to... | Use this command... |
|--|---|
| Display the health status of the system, which reflects the overall status of individual health monitors | <code>system health status show</code> |
| Display the health status of subsystems for which health monitoring is available | <code>system health subsystem show</code> |

Display the status of node connectivity

| If you want to... | Use this command... |
|--|---|
| Display details about connectivity from the node to the storage shelf, including port information, HBA port speed, I/O throughput, and the rate of I/O operations per second | <code>storage shelf show -connectivity</code> Use the <code>-instance</code> parameter to display detailed information about each shelf. |
| Display information about drives and array LUNs, including the usable space, shelf and bay numbers, and owning node name | <code>storage disk show</code> Use the <code>-instance</code> parameter to display detailed information about each drive. |
| Display detailed information about storage shelf ports, including port type, speed, and status | <code>storage port show</code> Use the <code>-instance</code> parameter to display detailed information about each adapter. |

Manage the discovery of cluster, storage, and management network switches

| If you want to... | Use this command.. (ONTAP 9.8 and later) | Use this command.. (ONTAP 9.7 and earlier) |
|--|--|--|
| Display the switches that the cluster monitors | <code>system switch ethernet show</code> | <code>system cluster-switch show</code> |

| If you want to... | Use this command.. (ONTAP 9.8 and later) | Use this command.. (ONTAP 9.7 and earlier) |
|--|--|---|
| Display the switches that the cluster currently monitors, including switches that you deleted (shown in the Reason column in the command output), and configuration information that you need for network access to the cluster and management network switches. This command is available at the advanced privilege level. | <code>system switch ethernet show-all</code> | <code>system cluster-switch show-all</code> |
| Configure discovery of an undiscovered switch | <code>system switch ethernet create</code> | <code>system cluster-switch create</code> |
| Modify information about a switch that the cluster monitors (for example, device name, IP address, SNMP version, and community string) | <code>system switch ethernet modify</code> | <code>system cluster-switch modify</code> |
| Disable monitoring of a switch | <code>system switch ethernet modify -disable-monitoring</code> | <code>system cluster-switch modify -disable-monitoring</code> |
| Disable discovery and monitoring of a switch and delete switch configuration information | <code>system switch ethernet delete</code> | <code>system cluster-switch delete</code> |
| Permanently remove the switch configuration information which is stored in the database (doing so reenables automatic discovery of the switch) | <code>system switch ethernet delete -force</code> | <code>system cluster-switch delete -force</code> |
| Enable automatic logging to send with AutoSupport messages. | <code>system switch ethernet log</code> | <code>system cluster-switch log</code> |

Respond to generated alerts

| If you want to... | Use this command... |
|--|---------------------------------------|
| Display information about generated alerts, such as the resource and node where the alert was triggered, and the alert's severity and probable cause | <code>system health alert show</code> |

| If you want to... | Use this command... |
|--|---|
| Display information about each generated alert | <code>system health alert show -instance</code> |
| Indicate that someone is working on an alert | <code>system health alert modify</code> |
| Acknowledge an alert | <code>system health alert modify -acknowledge</code> |
| Suppress a subsequent alert so that it does not affect the health status of a subsystem | <code>system health alert modify -suppress</code> |
| Delete an alert that was not automatically cleared | <code>system health alert delete</code> |
| Display information about the AutoSupport messages that alerts triggered within the last week, for example, to determine whether an alert triggered an AutoSupport message | <code>system health autosupport trigger history show</code> |

Configure future alerts

| If you want to... | Use this command... |
|--|---|
| Enable or disable the policy that controls whether a specific resource state raises a specific alert | <code>system health policy definition modify</code> |

Display information about how health monitoring is configured

| If you want to... | Use this command... |
|---|--|
| Display information about health monitors, such as their nodes, names, subsystems, and status | <code>system health config show</code> <div>  <p>Use the <code>-instance</code> parameter to display detailed information about each health monitor.</p> </div> |
| Display information about the alerts that a health monitor can potentially generate | <code>system health alert definition show</code> <div>  <p>Use the <code>-instance</code> parameter to display detailed information about each alert definition.</p> </div> |

| If you want to... | Use this command... |
|---|--|
| Display information about health monitor policies, which determine when alerts are raised | <pre data-bbox="820 163 1421 193">system health policy definition show</pre> <div data-bbox="850 317 902 369">  </div> <div data-bbox="964 245 1453 445"> <p>Use the <code>-instance</code> parameter to display detailed information about each policy. Use other parameters to filter the list of alerts—for example, by policy status (enabled or not), health monitor, alert, and so on.</p> </div> |

Display environmental information

Sensors help you monitor the environmental components of your system. The information you can display about environmental sensors include their type, name, state, value, and threshold warnings.

Step

1. To display information about environmental sensors, use the `system node environment sensors show` command.

Copyright information

Copyright © 2023 NetApp, Inc. All Rights Reserved. Printed in the U.S. No part of this document covered by copyright may be reproduced in any form or by any means—graphic, electronic, or mechanical, including photocopying, recording, taping, or storage in an electronic retrieval system—without prior written permission of the copyright owner.

Software derived from copyrighted NetApp material is subject to the following license and disclaimer:

THIS SOFTWARE IS PROVIDED BY NETAPP “AS IS” AND WITHOUT ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE, WHICH ARE HEREBY DISCLAIMED. IN NO EVENT SHALL NETAPP BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

NetApp reserves the right to change any products described herein at any time, and without notice. NetApp assumes no responsibility or liability arising from the use of products described herein, except as expressly agreed to in writing by NetApp. The use or purchase of this product does not convey a license under any patent rights, trademark rights, or any other intellectual property rights of NetApp.

The product described in this manual may be protected by one or more U.S. patents, foreign patents, or pending applications.

LIMITED RIGHTS LEGEND: Use, duplication, or disclosure by the government is subject to restrictions as set forth in subparagraph (b)(3) of the Rights in Technical Data -Noncommercial Items at DFARS 252.227-7013 (FEB 2014) and FAR 52.227-19 (DEC 2007).

Data contained herein pertains to a commercial product and/or commercial service (as defined in FAR 2.101) and is proprietary to NetApp, Inc. All NetApp technical data and computer software provided under this Agreement is commercial in nature and developed solely at private expense. The U.S. Government has a non-exclusive, non-transferrable, nonsublicensable, worldwide, limited irrevocable license to use the Data only in connection with and in support of the U.S. Government contract under which the Data was delivered. Except as provided herein, the Data may not be used, disclosed, reproduced, modified, performed, or displayed without the prior written approval of NetApp, Inc. United States Government license rights for the Department of Defense are limited to those rights identified in DFARS clause 252.227-7015(b) (FEB 2014).

Trademark information

NETAPP, the NETAPP logo, and the marks listed at <http://www.netapp.com/TM> are trademarks of NetApp, Inc. Other company and product names may be trademarks of their respective owners.